Service shutdown report

Generated: 2026-10-16 11:23:56
Target: ids (local)
Operation: **stop** · Recorded window: 26s · compose exit code: **0**

Resource usage recorded while the stack was taken down. Each dashed red vertical line marks a major event (a container created, started, going healthy or unhealthy, a one-shot job finishing, or a container stopping or being removed), numbered in a band above the graph; events within 1.5s of each other share one number, and the event key under the host graph says what each number covers. The largest host CPU peak is called out with the marker it lines up with. **Figure 1** shows how these services fit together – the per-container graphs below track each box in it.

Reading the metrics. **SD** (standard deviation) is the typical distance of a value from the mean. The **95th percentile** is the level usage stayed at or below 95% of the time; a **peak** far above it is a short burst rather than sustained load. **P(≥x)** is the share of the recording spent at or above *x*: the chance that a randomly picked moment was that busy. "How unusual the peak is" counts SDs above the mean; more than about 3 is a rare spike.

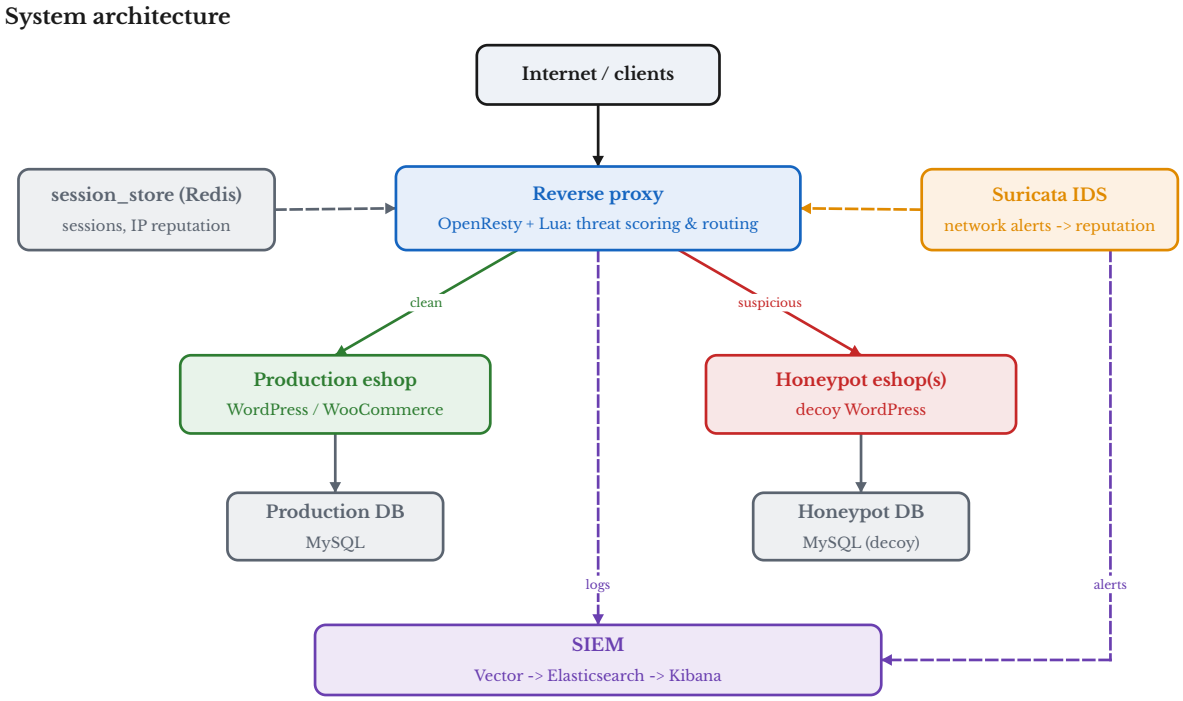
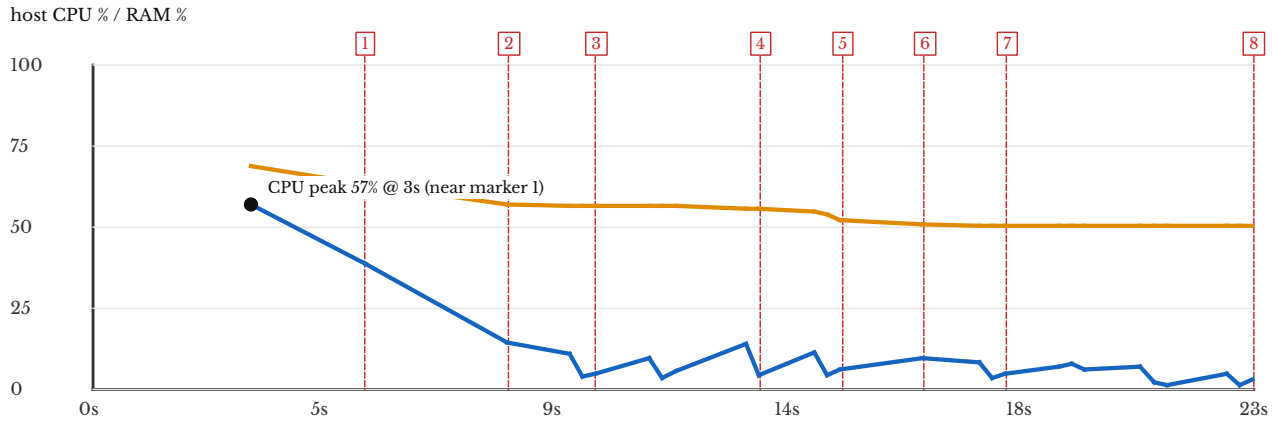


Figure 1. System architecture: the services whose CPU / memory the per-container charts below track as the stack is stopped.

1. Overall system (host)



Event key (marker numbers on all graphs)

Marker	Elapsed	What happened
1	5s	stopped: vector_outbound
2	8s	stopped: honeypot_eshop_1, reverse_proxy, suricata_ids
3	10s	stopped: honeypot_eshop_2, honeypot_eshop_3
4	13s	stopped: backup_service, honeypot_content_sync, pool_manager, log_capper
5	15s	finished: honeypot_database_3; stopped: production_eshop
6	16s	stopped: honeypot_database_1, honeypot_database_2; removed: honeypot_database_3
7	18s	stopped: production_database
8	23s	stopped: session_store

Probability metrics: host		
Host CPU, mean ± SD	**9.3% ± 11.9**	average over 27 samples; the SD is the typical swing around it
Host CPU, 95th percentile / peak	**31.2% / 57.0%**	95% of the time usage was at or below the first figure
Host CPU, how unusual the peak is	**4.0 SD above the mean**	more than ~3 SD is a rare spike rather than normal fluctuation
Host CPU, P(≥50%)	**3.7%**	chance a random moment of this recording was at least this busy
Host CPU, P(≥80%)	**0.0%**	chance a random moment of this recording was at least this busy
Host RAM used, mean ± SD	**53.8% ± 4.5**	average over 27 samples; the SD is the typical swing around it
Host RAM used, 95th percentile / peak	**61.1% / 68.7%**	95% of the time usage was at or below the first figure
Host RAM used, how unusual the peak is	**3.3 SD above the mean**	more than ~3 SD is a rare spike rather than normal fluctuation
Host RAM used, P(≥50%)	**100.0%**	chance a random moment of this recording was at least this busy
Host RAM used, P(≥80%)	**0.0%**	chance a random moment of this recording was at least this busy

2. Per-container

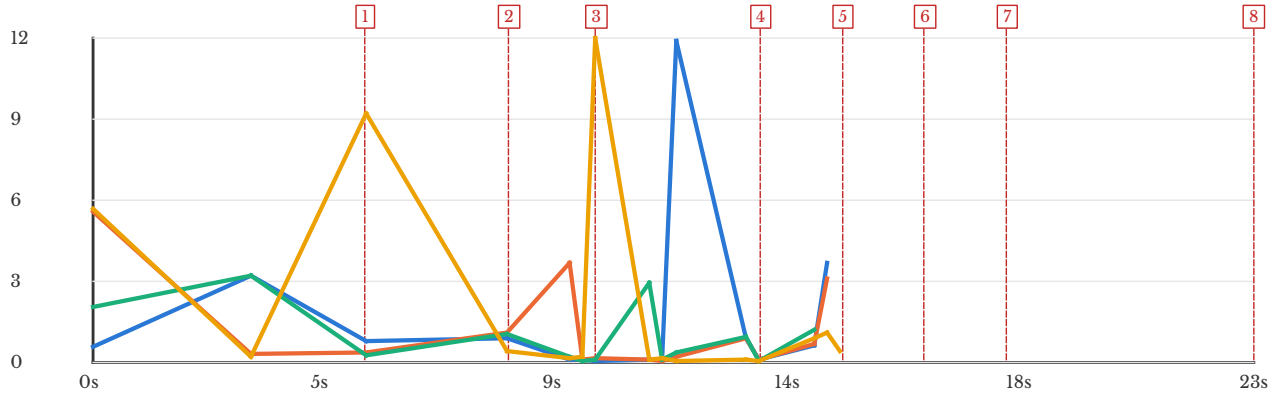
CPU and memory were read from the containers' cgroup files every 0.25 s.

Grouped by role, one CPU and one memory chart per group, so every line in a chart

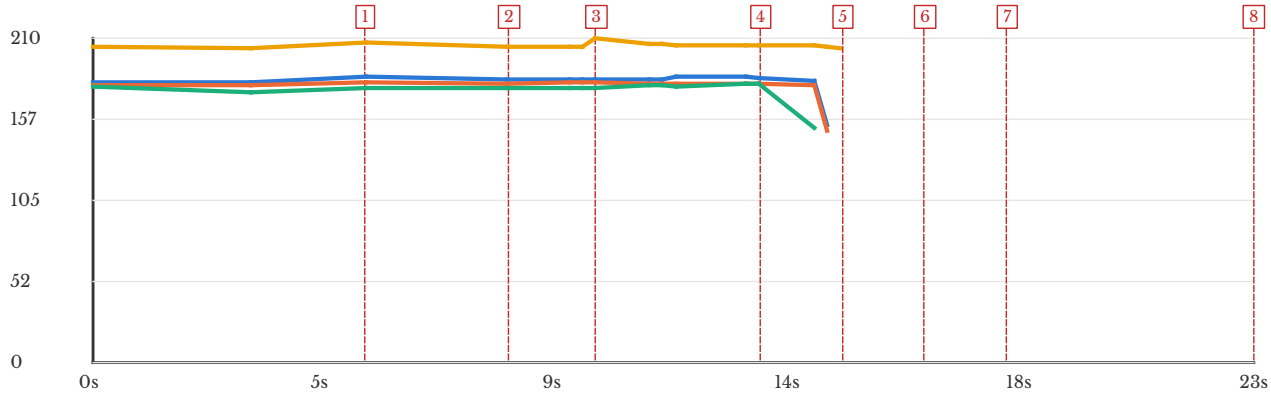
has a colour of its own.

Databases

Databases: container CPU %



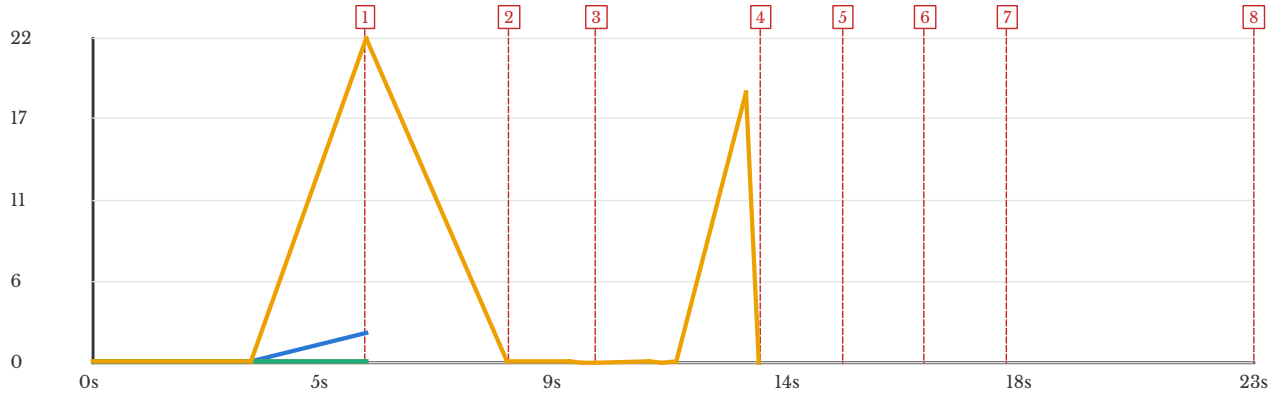
Databases: container memory (MiB)



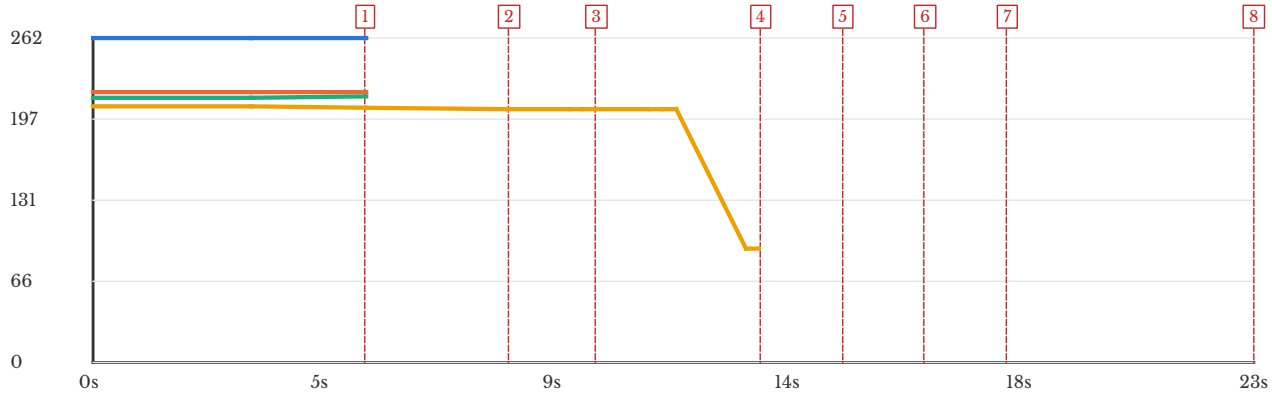
Legend: honeypot_database_1 honeypot_database_2 honeypot_database_3 production_database

Shops

Shops: container CPU %



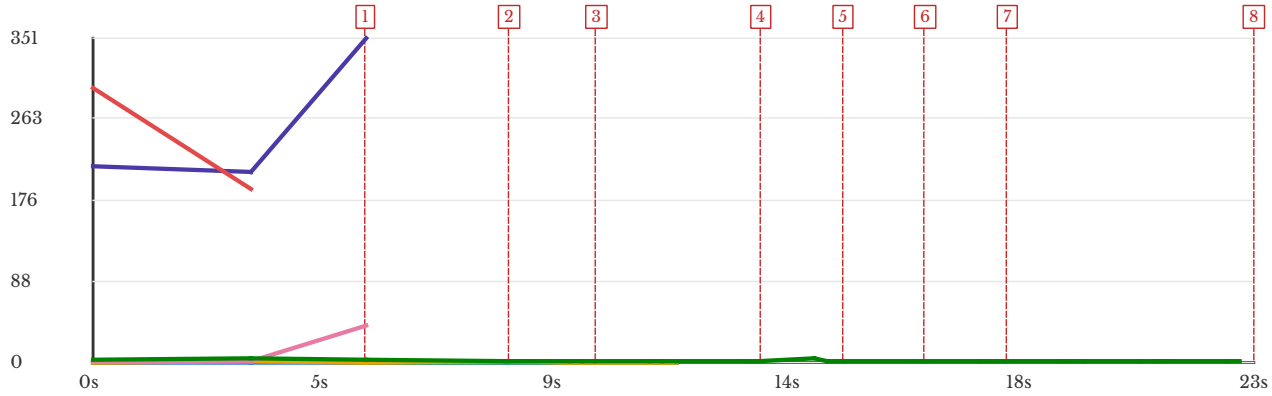
Shops: container memory (MiB)



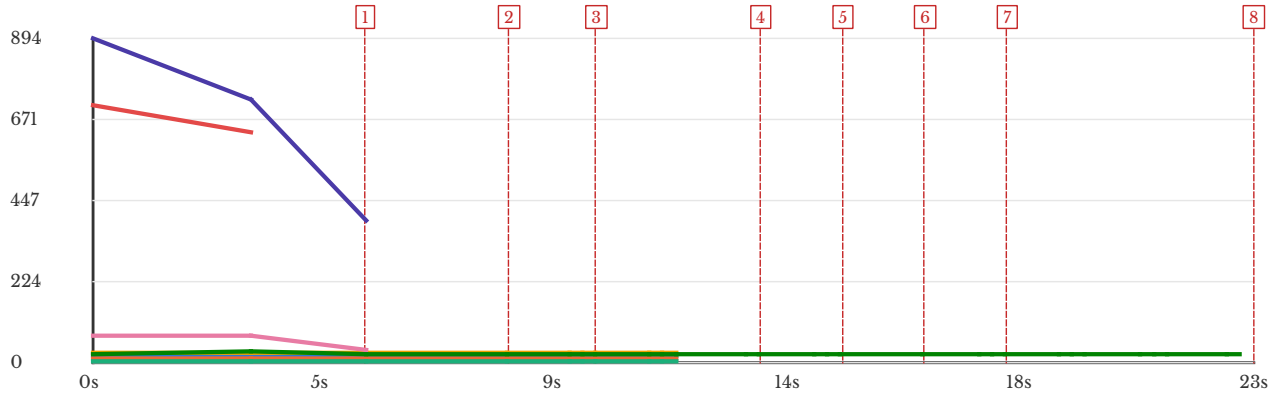
Legend: ■ honeypot_eshop_1 ■ honeypot_eshop_2 ■ honeypot_eshop_3 ■ production_eshop

Edge & monitoring

Edge & monitoring: container CPU %



Edge & monitoring: container memory (MiB)



Legend: ■ backup_service ■ honeypot_content_sync ■ log_capper ■ pool_manager ■ reverse_proxy ■ session_store ■ suricata_ids ■ vector_outbound

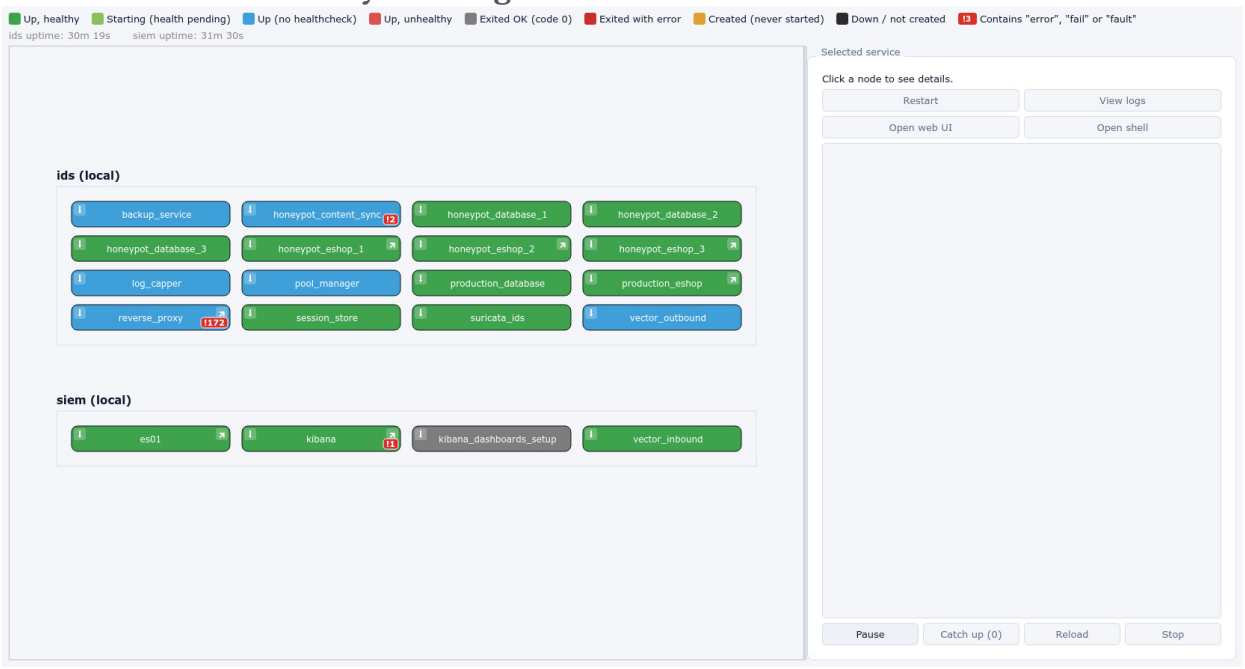
Probability metrics: per container					
Container	**CPU mean ± SD**	**CPU p95 / peak**	**P(CPU ≥ 50%)**	**Memory MiB, mean ± SD**	**Memory MiB, peak**
backup_service	0.0% ± 0.0	0.0% / 0.0%	0.0%	18 ± 0	18
honeypot_content_sync	0.0% ± 0.1	0.2% / 0.3%	0.0%	8 ± 0	8
honeypot_database_1	1.7% ± 3.3	6.7% / 12.2%	0.0%	181 ± 8	185
honeypot_datab	1.2% ± 1.8	4.5% / 5.7%	0.0%	178 ± 8	181

ase_2					
honeypot_database_3	1.0% ± 1.1	3.1% / 3.3%	0.0%	176 ± 7	180
honeypot_eshop_1	0.7% ± 1.1	1.8% / 2.0%	0.0%	262 ± 0	262
honeypot_eshop_2	0.0% ± 0.0	0.0% / 0.0%	0.0%	218 ± 0	218
honeypot_eshop_3	0.0% ± 0.0	0.0% / 0.0%	0.0%	214 ± 0	214
log_capper	0.1% ± 0.3	0.6% / 1.1%	0.0%	1 ± 0	1
pool_manager	0.2% ± 0.5	1.0% / 1.4%	0.0%	27 ± 0	27
production_database	2.1% ± 3.9	10.3% / 12.3%	0.0%	205 ± 2	210
production_eshop	3.4% ± 7.9	20.0% / 22.0%	0.0%	185 ± 44	207
reverse_proxy	13.2% ± 22.1	34.9% / 38.7%	0.0%	60 ± 23	73
session_store	0.9% ± 0.8	2.7% / 3.5%	0.0%	23 ± 1	29
suricata_ids	256.2% ± 82.3	337.3% / 351.2%	100.0%	671 ± 256	894
vector_outbound	241.6% ± 78.0	291.3% / 296.8%	100.0%	671 ± 55	710

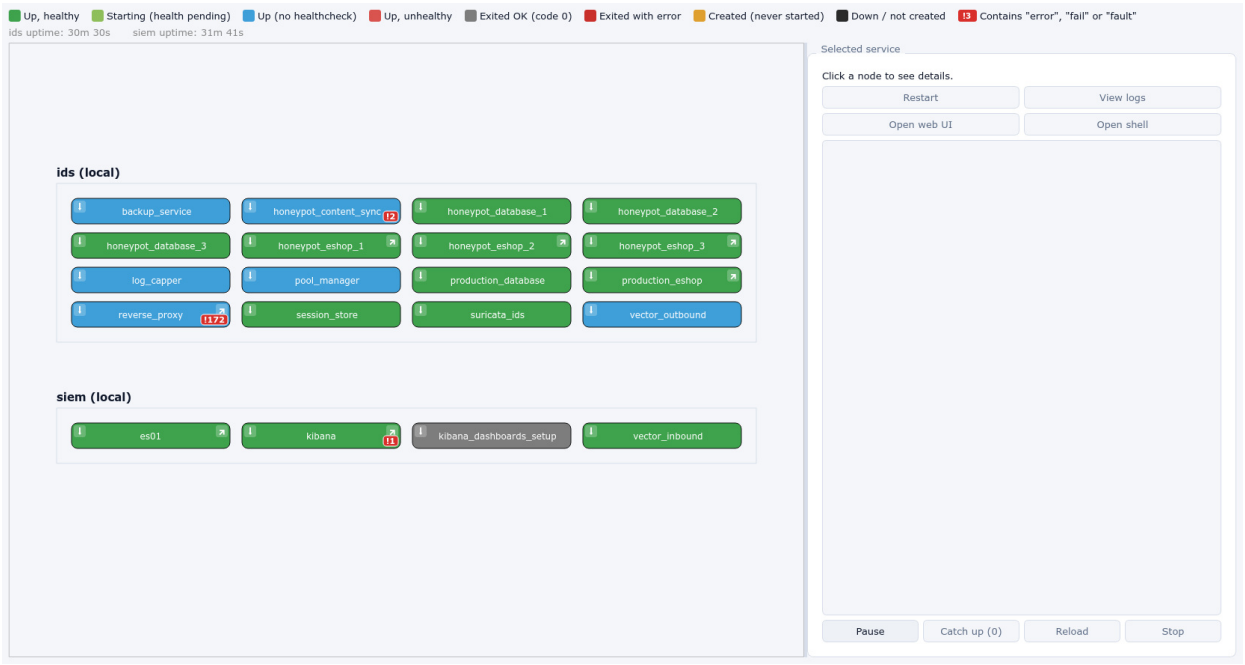
P(CPU ≥ 50%) is the share of the recording the container spent at half a core or more, i.e. the chance of catching it that busy at a random moment. A peak far above the p95 means a short burst (typically start-up work or a first healthcheck) rather than sustained load.

3. Global Visualization page as the stack changes state

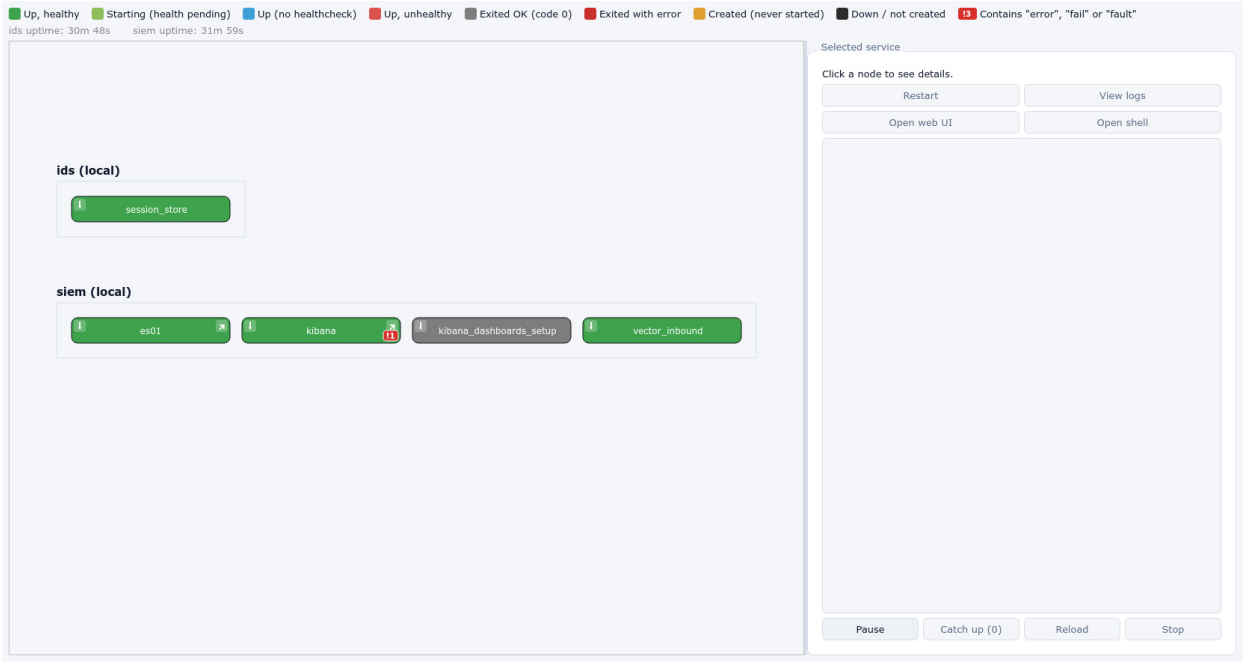
Before: stack idle / not yet changed



Stopping: containers going down



Stopped: down command finished



4. Event log

Marker	Elapsed	Event
1	5s	vector_outbound stopped
2	8s	honeypot_eshop_1 stopped
2	8s	reverse_proxy stopped
2	8s	suricata_ids stopped
3	10s	honeypot_eshop_2 stopped

3	10s	honeypot_eshop_3 stopped
4	13s	backup_service stopped
4	13s	honeypot_content_sync stopped
4	13s	pool_manager stopped
4	13s	log_capper stopped
5	15s	honeypot_database_3 finished
5	15s	production_eshop stopped
6	16s	honeypot_database_1 stopped
6	16s	honeypot_database_2 stopped
6	16s	honeypot_database_3 removed
7	18s	production_database stopped
8	23s	session_store stopped